Application of Equal Local Levels to Improve Q-Q Plot Testing Bands with R Package qqconf

Eric Weine
University of Chicago

Mary Sara McPeek
University of Chicago

Mark Abney
University of Chicago

Abstract

Quantile-Quantile plots (Q-Q plots) are often difficult to interpret because it is unclear how large the deviation from the theoretical distribution must be to indicate a lack of fit. Current packages and algorithms towards this end either do not ensure a robust Type I error rate, are too slow, or are under-powered to deviations in the tails of the distribution. In this paper, we present an efficient algorithm and accompanying R Package **qqconf** that computes simultaneous testing bands for Q-Q plots. These bands have a variety of desirable properties, including being fast to compute and being equally sensitive to deviations in all parts of the null distribution, including the tails.

Keywords: Q-Q plots, Equal Local Levels, Kolmolgorov-Smirnov, GWAS, Multiple Testing.

1. Introduction

Quantile-Quantile plots (Q-Q plots) are a common statistical tool used to judge if a sample comes from a specified distribution (Wilk and Gnanadesikan 1968). Despite their ubiquity, they are often difficult to interpret because it is challenging to determine if the magnitude of the deviation from the specified distribution is large enough to indicate a lack of fit as opposed to sampling variability. To make this determination, it is useful to put testing bands on a Q-Q plot.

A number of methods have been created towards this end. Naively, one could put pointwise testing bands onto Q-Q plots, which conduct an α level test on each order statistic of the sample. This method is clearly insufficient, as it ignores the multiple testing problem and thus provides no guarantee on the Type I error rate of testing the global null hypothesis that all of the data come from the specified null distribution. To appropriately deal with this multiple testing problem, the Komolgorov-Smirnov (KS)

statistic is commonly used to create joint testing bands (Kolmogorov 1941; Smirnov 1944). While this method ensures the correct Type I error by using the distribution of the KS statistic to create joint testing bands, this test suffers from very low power under a variety of reasonable alternatives since it has low sensitivity to deviation in the tails of the null distribution (Aldor-Noiman, Brown, Buja, Rolke, and Stine 2013; Berk and Jones 1979). A more promising approach would be to apply the equal local levels method (ELL) to create testing bands. This global testing method, originally introduced by Berk and Jones and later developed by Gontscharuk et al. as an alternative testing procedure to higher-criticism and KS, conducts a test with Type I error rate η on each order statistic of the sample such that the Type I error rate of the global test is some desired level α (Berk and Jones 1979; Gontscharuk, Landwehr, and Finner 2015; Gontscharuk, Landwehr, Finner *et al.* 2016; Gontscharuk and Finner 2017). Previously, Alder-Noiman applied this method to create testing bands for Q-Q plots for the normal distribution using simulation (Aldor-Noiman *et al.* 2013). Because this test conducts an η level test on each order statistic and thus does not bias the user towards deviations of the sample from any particular part of the null distribution, we suggest it is a sensible choice for a standard method in creating testing bands for Q-Q plots.

Currently, there is no available software that allows a user to quickly create plausible and robust testing bands for Q-Q plots on real data. While the base-R package **stats** provides functionality for creating a Q-Q plot to compare a sample against the normal distribution, it doesn't give the user the ability to create Q-Q plots for other distributions or to put testing bands on those plots. The package **qqplotr** provides a number of helpful additions to the base-R functionality, including the ability to create Q-Q plots for a variety of reference distributions and a robust method for testing bands on the normal distribution (Almeida, Loy, and Hofmann 2017). Specifically, for the normal distribution **qqplotr** implements Alder-Noiman's simulation based method (which is referred to in the package as tail sensitive (TS) testing bands). However, this method does not provide testing bands for non-normal reference distributions and is not able to efficiently create testing bands for large sample sizes because it uses simulation.

In what follows, we introduce the R package **qqconf**. With the aid of a novel recursive calculation developed below, **qqconf** can quickly provide ELL testing bands for comparing even very large samples to any reference distribution with a CDF implemented. In addition to these testing bands, **qqconf** provides a variety of plotting functionalities that allow the user to easily visualize where any deviation of the sample from the null distribution may occur. In section 2, we introduce the methods required for the efficient recursive calculation mentioned above. In section 3, we discuss the behavior of ELL bands in large sample sizes and an approximation to these bands that makes them very easy to calculate for even very large samples ($n > 5 \cdot 10^6$). In section 4, we introduce **qqconf** and demonstrate its functionality on a few multiple testing scenarios including including Type I error calibration and genetic studies. Finally, we conclude in section 5 with a discussion of some limitations of the ELL method and **qqconf** as well as further directions for creating testing bands on Q-Q plots.

2. Methods

2.1. Local Levels For Global Null Hypothesis Testing

Suppose we have observations

$$X_1, ..., X_n \overset{iid}{\sim} F,$$

and we are interested in conducting the following hypothesis test

$$H_0 : F = F_0 \text{ vs. } H_A : F \neq F_0$$

with level α, where all parameters of F_0 are known. One approach to this problem, referred to as "local levels," is to conduct separate hypothesis tests on each of the order statistics $X_{(i)}$, where each test has Type I error rate η_i (Gontscharuk *et al.* 2016). Then, we reject the global null hypothesis if at least one of the n tests results in a rejection. That is, we construct a set of intervals

$$(h_1, g_1), ..., (h_n, g_n),$$

and we reject H_0 if for any i

$$X_{(i)} \notin (h_i, g_i).$$

If we know the local level η_i, then

$$h_i = F_i^{-1}(\eta_i/2) \text{ and } g_i = F_i^{-1}(1 - \eta_i/2), \tag{1}$$

where F_i is the CDF of the i^{th} order statistic. Thus, the difficulty is in determining an appropriate vector of local levels $(\eta_1, ..., \eta_n)$ for which the resulting global level is α. In our case, we want to create testing bands that are "agnostic" to any alternative distribution. By this, we mean that we would like to design a test that applies equal scrutiny to each order statistic, and thus we set

$$\eta_1 = \eta_2 = ... = \eta_n = \eta. \tag{2}$$

We refer to this as the ELL method, as mentioned in the Introduction.

Given the sample size n and the desired level α we present the following algorithm for obtaining η in two steps:

1. Given a proposed testing region $(h_1, g_1), ..., (h_n, g_n)$, where $h_i < g_i$ for $i = 1, ..., n$, find the probability of $X_{(1)}, ..., X_{(n)}$ falling outside of this region under the null.

2. Using the algorithm described in step 1, conduct a binary search on the space of the local level η. This involves proposing a value of η, calculating the corresponding region $(h_1, g_1), ..., (h_n, g_n)$ using equation (1), and then using step 1 to check if the region achieves the desired global Type I error rate α.

We now describe these two steps in more detail:

Step 1: First, note that under the null hypothesis,

$$F_0(X_1), ..., F_0(X_n) \overset{iid}{\sim} U(0, 1).$$

Thus, without loss of generality, we can build the algorithm for step 1 assuming that

$$X_1, ..., X_n \overset{iid}{\sim} U(0, 1)$$

0	h1		h2		g1		g2		h3		g3		1
$u_1 = 0$		$u_2 = 1$		$u_3 = 2$		$u_4 = 2$		$u_5 = 2$		$u_6 = 3$			
$l_1 = 0$		$l_2 = 0$		$l_3 = 1$		$l_4 = 2$		$l_5 = 2$		$l_6 = 3$			

Figure 1: Example of how l_k and u_k, $k = 1, ..., 2n$, depend on the relative positions of $h_1, ..., h_n$ and $g_1, ..., g_n$, for a case with n = 3.

We are trying to calculate the following probability

$$\alpha = P_0\Big(\bigcup_{i=1}^{n} X_{(i)} \notin (h_i, g_i) \Big)$$

Where P_0 is the probability under the null hypothesis. It is easier to calculate this as

$$\alpha = 1 - P_0\Big(\bigcap_{i=1}^{n} X_{(i)} \in (h_i, g_i) \Big)$$

To calculate this value, it is easier to transform the events over order statistics to multinomial events. To do this, we introduce the following notation:

Let $b_1, ..., b_{2n}$ be the sorted values of $h_1, ..., h_n, g_1, ..., g_n$ in ascending order. We also define $b_0 = 0$ and $b_{2n+1} = 1$. We will divide the interval (b_0, b_{2n+1}) into $2n + 1$ bins, where bin 1 is (b_0, b_1), bin 2 is (b_1, b_2), ..., and bin $2n+1$ is (b_{2n}, b_{2n+1}). Note that because the uniform distribution is continuous, the probabilities associated with the endpoints of the bins are 0, so it does not matter if the left and right boundaries of these intervals are open or closed. Let $N_j = \sum_{i=1}^{n} \mathbb{1}\Big(X_i \in (b_{j-1}, b_j) \Big)$ denote the random variable that counts the number of X's falling into bin j, for $1 \leq j \leq 2n + 1$, and let $S_k = \sum_{j=1}^{k} N_j$ be the *kth* partial sum of the *N*'s, for $1 \leq k \leq 2n + 1$. The key to the algorithm is that the following two events are the same:

$$\{X_{(i)} \in (h_i, g_i) \text{ for } i = 1, ..., n\} = \{l_k \leq S_k \leq u_k \text{ for } k = 1, 2, ..., 2n\},$$

where

$$u_k = \begin{cases} 0, & \text{if } k = 1 \\ \sum_{i=1}^{k-1} \mathbb{1}\big(b_i \in \{h_1, ..., h_n\}\big), & \text{otherwise} \end{cases}$$

$$l_k = \sum_{i=1}^{k} \mathbb{1}\big(b_i \in \{g_1, ..., g_n\}\big)$$

(See Figure 1 above for intuition). Finally, we define

$$c_w^{(k)} = P_0(S_k = w \text{ and } l_j \leq S_j \leq u_j \text{ for } j = 1, ..., k-1), \text{for } k = 1, ..., 2n \text{ and } w = 1, ..., n.$$

Our basic approach to step 1 is a recursive calculation of these $c_w^{(k)}$ values. Then, the global level α could be computed as $1 - c_n^{(2n)}$ in the general case. However, for the special case in which $(h_1, g_1), ..., (h_n, g_n)$ are derived from equal local levels, i.e. equation (2) holds, then as a result of the symmetry in the problem, we need only calculate $c_w^{(k)}$ for $k = 1, ..., n + 1$ instead of $k = 1, ..., 2n$. This shortcut is implemented below in the

pseudocode for step 1. The algorithm calculates the probability of a valid allocation (under the null hypothesis) of points in bins by looping over each bin and making a recursive calculation.

Algorithm 1 Calculate Type I error α from proposed rejection region

---

Input: Vector of lower bound values $(h_1, ..., h_n)$, vector of upper bound values $(g_1, ..., g_n)$, where we require $h_i < g_i$ for $i = 1, ..., n$.

$get\_level\_from\_bounds\_two\_sided(h_1, ..., h_n, g_1, ..., g_n)$

1: $b_1, ..., b_{2n} \leftarrow \mathbf{sort}(h_1, ..., h_n, g_1, ..., g_n)$
2: $c_0^{(1)} \leftarrow (1 - b_1)^n$
3: $l_k \leftarrow 0$
4: $u_k \leftarrow 0$
5: **for** $k = 2, ..., n + 1$ **do**
6:     **if** $b_{k-1} \in \{h_1, ..., h_n\}$ **then**
7:         $u_k \leftarrow u_{k-1} + 1$
8:     **end if**
9:     **if** $b_k \in \{g_1, ..., g_n\}$ **then**
10:        $l_k \leftarrow l_{k-1} + 1$
11:     **end if**
12:     **for** $j = l_k, ..., u_k$ **do**
13:        $c_j^{(k)} \leftarrow 0$
14:        **for** $m = l_{k-1}, ..., min(u_{k-1}, j)$ **do**
15:           $c_j^{(k)} \leftarrow c_j^{(k)} + c_m^{(k-1)} * dbinom(x = j - m, size = n - m, prob = \frac{(b_k - b_{k-1})}{(1 - b_{k-1})})$
16:        **end for**
17:     **end for**
18: **end for**
19: $1 - \alpha \leftarrow 0$
20: **for** $l = l_n, ..., u_n$ **do**
21:     $1 - \alpha \leftarrow 1 - \alpha + \frac{c_l^{(n)} * c_{n-l}^{(n+1)}}{dbinom(x=l, size=n, prob=b_n)}$
22: **end for**
23: **return** α

end

Now, we derive the algorithm. **Initialization:** (Lines 2 through 4 of Algorithm 1)

$$
\begin{aligned}
c_0^{(1)} &= P_0(S_1 = 0) \\
&= P_0\Big(\bigcap_{i=1}^{n}\{X_i \in (b_1, 1)\}\Big) \\
&= \prod_{i=1}^{n} P_0\Big(X_i \in (b_1, 1)\Big) && \text{(by independence)} \\
&= (1 - b_1)^n && \text{(Since each } X_i \sim U(0,1) \text{ under the null)}
\end{aligned}
$$

Also, we initialize $l_1 = 0$ and $u_1 = 0$ because the only allowable partial sum is 0 in the

first bin. **Recursion:** (Line 15 of Algorithm 1)

$$
\begin{aligned}
c_j^{(k)} &= P_0(S_k = j \text{ and } l_q \le S_q \le u_q \text{ for } q = 1, ..., k-1) \\
&= P_0\Big(\bigcup_{m=l_{k-1}}^{min(j,u_{k-1})} \{S_{k-1} = m \text{ and } N_k = j-m \text{ and } l_q \le S_q \le u_q \text{ for } q = 1, ..., k-2\}\Big) \\
&= \sum_{m=l_{k-1}}^{min(j,u_{k-1})} P_0(S_{k-1} = m \text{ and } N_k = j-m \text{ and } l_q \le S_q \le u_q \text{ for } q = 1, ..., k-2) \\
&= \sum_{m=l_{k-1}}^{min(j,u_{k-1})} P_0(S_{k-1} = m \text{ and } l_q \le S_q \le u_q \text{ for } q = 1, ..., k-2) * \\
&\qquad P_0(N_k = j-m | S_{k-1} = m \text{ and } l_q \le S_q \le u_q \text{ for } q = 1, ..., k-2) \\
&= \sum_{m=l_{k-1}}^{min(j,u_{k-1})} c_m^{(k-1)} * P_0(N_k = j-m | S_{k-1} = m) \\
&= \sum_{m=l_{k-1}}^{min(j,u_{k-1})} c_m^{(k-1)} * P(B = j-m), \text{where } B \sim Binomial\left(n-m, \frac{b_k - b_{k-1}}{1 - b_{k-1}}\right),
\end{aligned}
$$

which is easily computed.

As noted above, in the general case, we could use this recursion to obtain $c_n^{(2n)}$, and then compute $\alpha = 1 - c_n^{(2n)}$. However, in the special case of equal local levels, we can use symmetry to speed the calculation by only computing the $c_w^{(k)}$ values for $k = 1, ..., n+1$ and $w = 1, ..., n$, and performing the calculation on line 21 of Algorithm 1, as we now demonstrate. First, we define the following values:

$$
\tilde{u}_k = \begin{cases} 0, & \text{if } k = 2n \\ \sum_{i=k+1}^{2n} \mathbb{1}(b_i \in \{g_1, ..., g_n\}), & \text{otherwise} \end{cases}
$$

$$
\tilde{l}_k = \sum_{i=k}^{2n} \mathbb{1}(b_i \in \{h_1, ..., h_n\})
$$

$$
T_k = \sum_{j=k+1}^{2n+1} N_j = n - S_k
$$

Now, we make the following observations:

(1) With equal local levels, $g_i = 1 - h_{n+1-i}$ for $i = 1, ..., n$ since $F^{-1}_{Beta(i,n+1-i)}(1 - \frac{\eta}{2}) = 1 - F^{-1}_{Beta(n+1-i,i)}(\frac{\eta}{2})$.

(2) $b_k = 1 - b_{2n+1-k}$ for $k = 1, ..., 2n$ by (1).

(3) $u_k = \tilde{u}_{2n+1-k}$ and $l_k = \tilde{l}_{2n+1-k}$, for $k = 1, ..., 2n$ by (1) and (2).

(4) The random vector $(N_1, ..., N_k)$ has the same distribution as $(N_{2n+1}, ..., N_{2n+2-k})$ for $k = 1, ..., 2n+1$. This follows from the fact that the vector $(X_1, ..., X_n)$ has the same distribution as $(1 - X_1, ..., 1 - X_n)$ (since each X_i is independent uniform) and (2).

(5) The random vector $(S_1, ..., S_k)$ has the same distribution as $(T_{2n}, ..., T_{2n+1-k})$ for $k = 1, ..., 2n$. This follows from (4).

(6)

$$c_j^{(k)} = P_0(S_k = j \text{ and } l_q \leq S_q \leq u_q \text{ for } q = 1, ..., k-1)$$
$$= P_0(T_{2n+1-k} = j \text{ and } \tilde{l}_r \leq T_r \leq \tilde{u}_r \text{ for } r = 2n+2-k, ..., 2n).$$

This follows from (3) and (5).

(7) Conditional on S_k, the random vector $(X_{(S_k+1)}, ..., X_{(n)})$ is distributed as the order statistics of $n - S_k$ $U(b_k, 1)$ independent random variables.

(8) The random vector $(S_1, ..., S_r)$ and the random vector $(T_r, ..., T_n)$ are independent conditional on S_r. This follows directly from (7).

Combining the above results, we can write

$$\{X_{(i)} \in (h_i, g_i) \text{ for } i = 1, ..., n\} = \{l_k \leq S_k \leq u_k \text{ for } k = 1, ..., 2n\}$$

Also, observe that for any $2 \leq r \leq 2n - 1$, we have

$$\{l_k \leq S_k \leq u_k \text{ for } k = 1, ..., 2n\} =$$

$$\{l_k \leq S_k \leq u_k \text{ for } k = 1, ..., r \text{ and } T_r = n - S_r \text{ and } \tilde{l}_k \leq T_q \leq \tilde{u}_k \text{ for } q = r, ..., 2n\}$$

Thus, we can write

$$\sum_{j=l_r}^{u_r} P(S_r = j \text{ and } l_k \leq S_k \leq u_k \text{ for } k = 1, ..., r-1)$$

$$\cdot I(\tilde{l}_r \leq n - j \leq \tilde{u}_r) \cdot P(\tilde{l}_q \leq T_q \leq \tilde{u}_r \text{ for } q = r+1, ..., 2n | T_r = n - j)$$

$$= \sum_{j=l_r}^{u_r} c_j^{(r)} \cdot I(\tilde{l}_r \leq n - j \leq \tilde{u}_r) \cdot P(\tilde{l}_q \leq T_q \leq \tilde{u}_r \text{ for } q = r+1, ..., 2n | T_r = n - j)$$

$$= \sum_{j=l_r}^{u_r} c_j^{(r)} \cdot I(\tilde{l}_r \leq n - j \leq \tilde{u}_r) \cdot \frac{P(\tilde{l}_q \leq T_q \leq \tilde{u}_r \text{ for } q = r+1, ..., 2n \text{ and } T_r = n - j)}{T_r = n - j}$$

$$= \sum_{j=l_r}^{u_r} c_j^{(r)} \cdot \frac{c_{n-j}^{(2n+1-r)}}{\binom{n}{j} b_r^j (1 - b_r)^{n-j}}$$

Now, if we let $r = n$ above, then we get

$$P_0\Big(\bigcap_{i=1}^n X_{(i)} \in (h_i, g_i)\Big) = \sum_{j=l_n}^{u_n} c_j^{(n)} \cdot \frac{c_{n-j}^{(n+1)}}{\binom{n}{j} b_n^j (1 - b_n)^{n-j}}$$

which is line 21 in the psuedo-code above.

Empirically, it seems that our algorithm has computational complexity $O(n^2)$. For a

dense grid of values of n between 10 and 50,000, the number of recursive steps required (line 15 in algorithm 1 above) scales with n^2. More precisely, it appears that the number of recursive steps required is very close to $8n^2$. While each recursive step itself requires calculating a binomial probability which is $O(n)$ due to the calculation of the binomial coefficient multiplied by a quantity with powers as large as n, these calculations can be memoized with $O(n^2)$ cost. Thus, these binomial probability calculations do not change the big O complexity of the algorithm.

Step 2: This step is done with a simple binary search over η. We begin by setting $\eta_{upper} = \frac{-log(1-\alpha)}{(2*log(log(n))*log(n))}$ since this is an upper bound on the local level as shown by Gontscharuk et al., and we set $\eta_{lower} = \frac{\alpha}{n}$, as this is the lower bound given by the Bonferroni correction. The pseudocode is shown below:

Algorithm 2 Calculate testing bounds from global level α and sample size n

---

Input: Local level α, sample size n, tolerance ϵ

$get\_bounds\_two\_sided(\alpha, n, \epsilon)$

  1: $\eta_{upper} \leftarrow \frac{-log(1-\alpha)}{(2*log(log(n))*log(n))}$

  2: $\eta_{lower} \leftarrow \frac{\alpha}{n}$

  3: $\alpha_{mid} \leftarrow \infty$

  4: **while** $\frac{\alpha_{mid}-\alpha}{\alpha} > \epsilon$ **do**

  5:      $\eta_{mid} \leftarrow \frac{\eta_{upper}+\eta_{lower}}{2}$

  6:      $h_1, ..., h_n \leftarrow qbeta(p = \frac{\eta_{mid}}{2}, \text{shape1} = c(1{:}n), \text{shape2} = c(n{:}1))$

  7:      $g_1, ..., g_n \leftarrow qbeta(p = 1 - \frac{\eta_{mid}}{2}, \text{shape1} = c(1{:}n), \text{shape2} = c(n{:}1))$

  8:      $\alpha_{mid} \leftarrow get\_level\_from\_bounds\_two\_sided(h_1, ..., h_n, g_1, ..., g_n)$

  9: **end while**

10: **return** $h_1, ..., h_n, g_1, ..., g_n$

end

---

3. Local Level Approximations in Large Samples

For sufficiently large values of the number of tests n, it can be expedient to apply an accurate asymptotic approximation of the local level η corresponding to global level α in place of exact computation. Previously, Gontscharuk and Finner showed that the asymptotic local level for global level α and number of tests n is

$$\eta_{asymp} = \frac{log(1-\alpha)}{2log(log(n))log(n)}$$

(Gontscharuk and Finner 2017). However, as they note, this approximation gives poor performance for n even as large as 10^4. To improve this approximation, they propose to add a smaller order correction term, resulting in an approximation of the form

$$\eta_{approx} = \frac{log(1-\alpha)}{2log(log(n))log(n)} \left[1 - c_\alpha \frac{log(log(log(n)))}{log(log(n))} \right], \tag{3}$$

where c_α is chosen empirically. For the values $\alpha = .01, .05,$ and $.1$ they chose $c_\alpha = 1.6, 1.3,$ and 1.1, respectively. To select these c_α values, the authors calculated the values of η to high precision on a grid of values up until $n = 10,000$. To further test

these approximations, we decided to calculate the values of η with high precision on a grid of values up to $n = 500,000$ for $\alpha = .01, .05$. Based on our evaluation, we found that $c_\alpha = 1.3$ is satisfactory for $\alpha = .05$, but that $c_\alpha = 1.6$ for $\alpha = .01$ was not sufficiently accurate for our purposes (see Figure 2). This approximation begins anti-conservative for small values of n, becomes very accurate for values of n around 2,500, and then becomes and remains substantially conservative for even very large values of n.

In addition to implementing algorithms 1 and 2, which, in principle, can be applied for any n and α, our package offers a faster approximate approach specifically for $\alpha = .01$ and $.05$. For smaller values of n, up to $100,000$, we have pre-calculated values of η on a dense grid of values of n. If the user inputs a value of n less than or equal to $100,000$, we either return back the pre-computed value of η if n happens to be a grid point, or we use linear interpolation if the value of n is between grid points, which leads to a highly accurate approximation. If the user inputs a value of n greater than $100,000$, we use the asymptotic approximation given in (3), but with $c_\alpha = 1.591$ and 1.3 for $\alpha = .01$ and $.05$, respectively. We found that $c_\alpha = 1.591$ led to much better performance for $\alpha = .01$ for large values of n, as can be seen in Figure 2. Our package also implements the approximation given in equation (3) for $\alpha = .1$ with $c_\alpha = 1.1$, but we did not pre-compute any grid for these values due to limited computational resources.

4. Examples

As noted above, the algorithm presented is suitable for the case when all of the parameters of F_0 are known, and all $X_1, ..., X_n$ are independent. While these assumptions may seem restrictive, the methods above are still useful in a variety of scenarios.

One of the main advantages of the local levels method compared to other global testing approaches, is that it can easily be used to put testing bands onto Q-Q plots by simply graphing each (h_i, g_i) interval. This allows us to examine how a dataset might deviate from some null distribution much better than simply applying a test that yields a binary conclusion. Below, we present a few examples where a Q-Q plot is useful, and where the local levels test seems ideal for assessing deviation from a global null hypothesis.

4.1. Examining the P-value Distribution for Testing Procedures

Suppose we have devised a new testing procedure to test a null hypothesis H_0 with test statistic T, where we also specify a particular method to calculate or approximate p-values. In such a situation it is important to perform some simulations under the null hypothesis and check that the resulting p-value distribution is approximately uniform in the simulation experiment.

Typically, the verification of Type I error rate is done using the following procedure:

(1) Generate n simulated datasets under H_0, and calculate T for each siulated dataset to obtain $T_1, ..., T_n$.

(2) Select a value of α, and for each of $T_1, ..., T_n$, determine whether the null hypothesis is rejected at level α. Let N_α be the observed number of the n tests that are rejected

Figure 2: Relative Error of Local Level vs. Number of Tests for η_{approx} for $\alpha = .01$ with $c_\alpha = 1.6$ (top) $c_\alpha = 1.591$ (bottom).

at level α.

(3) Let α^* denote the true probability of rejection under the above procedure. Test the null hypothesis $H_0 : \alpha^* = \alpha$ by applying, e.g., a z-test of proportions or an exact binomial test to the data N_α.

While the above procedure provides reliable information about the Type I error calibration for one level of α, it provides little information about the global calibration of p-values. Instead, we suggest the following procedure:

(1) As above.

(2) For each T_i, calculate the corresponding p-value, p_i, to obtain $p_1, ..., p_n$.

(3) Make a q-q plot comparing $p_1, ..., p_n$ to a $U(0, 1)$ distribution, and apply the local levels procedure to create simultaneous testing bands for the null hypothesis that $p_1, ..., p_n \overset{iid}{\sim} U(0, 1)$.

This allows us to easily visualize the global calibration of the p-values with just one graph and diagnose any issues if they exist. In Step 3, one could use many different testing bands. However, in the calibration of p-values, we typically don't have the expectation that our p-values would be more likely to deviate from uniform in any particular region, and so it makes sense to use the local levels test because it is agnostic to the space of alternative distribution. Moreover, since it is generally most concerning if small p-values are not calibrated (i.e. those in the lower tail of the uniform distribution), the local levels test is preferable to the standard KS test because it is much more sensitive in the tails (Aldor-Noiman 2013).

Chi-Square Test for Independence in a 2 X 2 Table

We apply this approach to assess the calibration of p-values from the Pearson Chi-Square test for independence in a 2×2 table. A well-known rule of thumb is that the chi-square test is appropriate as long as the expected cell count in each cell under the null hypothesis is at least 5. We fix the cell probabilities under the null hypothesis and consider two different cases: in scenario 1, the sample size is $s = 200$ and the rule of thumb holds, and scenario 2, the sample size is only $s = 20$ and the rule of thumb does not hold. We use the local levels approach to generate simultaneous testing bands to assess the calibration of the p-values from the Pearson Chi-Square test for these two scenarios.

More specifically, in each scenario, we randomly generate $n = 1000$ 2×2 tables under the null hypothesis, where each table contains s observations, with $s = 200$ in scenario 1 and $s = 20$ in scenario 2. For each table, the s observations are i.i.d. with probability $q_{i,j}$ of falling in cell (i, j), for $i = 0, 1, j = 0, 1$, where $q_{1,1} = a * b$, $q_{1,0} = a * (1 - b)$, $q_{0,1} = (1 - a) * b$, and $q_{0,0} = (1 - a) * (1 - b)$, with a=.15 and b=.4. For each table, let $X_{i,j}$ denote the observed count in cell (i, j). If any table has $X_{0,0} + X_{0,1} = 0$ or $= s$, we discard the table and draw a new one, because that would imply that one of the rows of the table is empty, in which case the Pearson Chi-Squared test statistic is not defined. Similarly, if any table has $X_{0,0} + X_{1,0} = 0$ or $= s$, we discard the table and draw a new

one. The code to do this is shown below:

```
R> set.seed(8675309)
R> # Generate contingency tables for variables as defined above
R> generate_contingency_tables <- function(n, s, a, b) {
+
+     cell_probs <- c(a * b, a * (1 - b), b * (1 - a), (1 - b) * (1 - a))
+     num_sims <- 0
+     sim_mat <- matrix(NA, nrow = n, ncol = 4)
+
+     # Generate n simulations and append the contingency tables to a matrix
+     while (num_sims < n) {
+
+       proposed_sim <- t(rmultinom(n = 1, size = s, prob = cell_probs))
+
+       # Check to make sure that the chi-square statistic will be defined
+       if (proposed_sim[, 1] + proposed_sim[, 2] != 0 &&
+           proposed_sim[, 1] + proposed_sim[, 3] != 0 &&
+           proposed_sim[, 3] + proposed_sim[, 4] != 0 &&
+           proposed_sim[, 4] + proposed_sim[, 2]) {
+
+         num_sims <- num_sims + 1
+         sim_mat[num_sims, ] <- proposed_sim
+
+       }
+
+     }
+
+     return(sim_mat)
+
+ }
R> scenario_1 <- generate_contingency_tables(n = 1000, s = 200, a = .15, b = .4)
R> scenario_2 <- generate_contingency_tables(n = 1000, s = 20, a = .15, b = .4)
```

For each of the tables in the resulting sample, a Pearson chi-square test statistic is calculated as $T = \sum_{i=0}^{1} \sum_{j=0}^{1} \frac{(X_{i,j} - q_{i,j})^2}{q_{i,j}}$, where $X_{i,j}$ is the observed count in cell (i, j). For each scenario, this results in $n = 1000$ test statistics, $T_1, \ldots, T_n$, one for each table. From these, we obtain $n = 1000$ p-values, $p_1, \ldots, p_n$ by applying the χ_1^2 approximation, i.e., $p_i = 1 - F(T_i)$ for $i = 1, \ldots, n$, where F is taken to be the c.d.f. of the χ_1^2 distribution. The code to do this is shown below:

```
R> # Compute Pearson Chi-Square Test p-values for Set of Contingency Tables
R> get_p_values_from_c_tables <- function(c_tables, s) {
+
+     chisq_numerator <- (
+       (c_tables[, 1] * c_tables[, 4] - c_tables[, 2] * c_tables[, 3]) / s
+     ) ^ 2
+     chisq_denominator <- (c_tables[, 1] + c_tables[, 2]) *
+       (c_tables[, 1] + c_tables[, 3]) *
```

```
+       (c_tables[, 3] + c_tables[, 4]) *
+       (c_tables[, 2] + c_tables[, 4]) /
+       (s ^ 3)
+     chisq_stats_vec <- chisq_numerator / chisq_denominator
+     p_vals <- pchisq(q = chisq_stats_vec, df = 1, lower.tail = FALSE)
+     return(p_vals)
+
+   }
R> pvals_scenario_1 <- get_p_values_from_c_tables(scenario_1, s = 200)
R> pvals_scenario_2 <- get_p_values_from_c_tables(scenario_2, s = 20)
```

Figure 3 shows the resulting Q-Q plots for scenarios 1 (in blue) and 2 (in red), where the 45° line is shown as well as the testing bands obtained from the equal local levels procedure for testing, at global level .05, the null hypothesis that $p_1, \ldots, p_n$ have the same distribution as n i.i.d. draws from Unif(0,1).

```
R> library(qqconf)
R> n <- 1000
R> bounds <- get_bounds_two_sided(alpha = .05, n = 1000)
R> # Plot 3: Q-Q plot of P-values from Chi-Squared Test
R> # for Independence in a 2 x 2 Table
R>
R>
R> qq_conf_plot(
+     obs = pvals_scenario_2,
+     distribution = qunif,
+     method = "ell",
+     alpha = .05,
+     points_params = list(col="red", type="l")
+   )
R> # Add the Q-Q plot for scenario 1.
R> lines(c(1:n) / (n + 1), sort(pvals_scenario_1), col = "blue")
R> legend(
+     "topleft",
+     legend = c("s=20","s=200"),
+     col = c("red","blue"),
+     lty = 1
+   )
```

Figure 4 shows the same plots where the axes are transformed to be on the $-\log_{10}$ scale to focus the plot on small p-values. (Note that in Figure 4, small p-values are to the top and right of the plot, so a curve that is too low is conservative, and too high is anti-conservative.)

```
R> # Plot 4: Q-Q plot, on the -log10 Scale, of P-values
R> # from Chi-Squared Test for Independence in a 2 x 2 Table
R> n <- 1000
R> qq_conf_plot(
+     obs = pvals_scenario_2,
```

Figure 3: QQ plot for Scenario 1 and Scenario 2 With Standard Axes.

Figure 4: QQ plot for Scenario 1 and Scenario 2 With Axes On Log Scale.

```
+     distribution = qunif,
+     method = "ell",
+     alpha = .05,
+     points_params = list(col="red", type="l"),
+     log10 = TRUE
+ )
R> # Add the Q-Q plot for scenario 1.
R> lines(-log10(c(1:n) / (n + 1)), -log10(sort(pvals_scenario_1)), col = "blue")
R> legend(
+     "topleft",
+     legend = c("s=20","s=200"),
+     col = c("red","blue"),
+     lty = 1
+ )
```

From these plots, it can be seen that in scenario 1, when $s = 200$ and the smallest expected cell count is 12, there is no significant deviation of the p-values from i.i.d. Unif(0,1) under the null hypothesis. In contrast, in scenario 2, when $s = 20$ and the smallest expected cell count is 1.2, the χ_1^2 asymptotic distribution is not an accurate approximation to the sampling distribution of T. As a result, we can see in Figures 3

and 4 that the p-values differ significantly from i.i.d. Unif(0,1) under the null hypothesis, with small p-values tending to be overly conservative, while the larger p-values tend to be anti-conservative.

4.2. Testing for Normality of Regression Residuals

When performing an ordinary least squares regression (OLS), it is common to assume that the error terms are idependentally drawn from a normal distribution. Specifically, suppose that we have n observations $\{\mathbf{x_i}, y_i\}_{i=1}^n$, where each y_i is a scalar and each $\mathbf{x_i}$ is a $1xp$ vector of regressors. Then, we model each y_i as

$$y_i = \mathbf{x_i}^T \boldsymbol{\beta} + \epsilon_i$$

where $\boldsymbol{\beta}$ is a $1xp$ vector of regression parameters to be estimated and we assume

$$\epsilon_1, ..., \epsilon_n \stackrel{iid}{\sim} N(0, \sigma^2). \tag{4}$$

Once $\boldsymbol{\beta}$ and σ^2 are estimated (typically via maximum liklihood), for each y_i we define the residual as

$$r_i = \mathbf{x_i}^T \hat{\boldsymbol{\beta}} - y_i,$$

where $\hat{\boldsymbol{\beta}}$ is our estimate of $\boldsymbol{\beta}$. Then, we would like a Q-Q plot of the residuals with a Normal reference distribution to aid in testing assumption (4) above. Strictly speaking, even if assumption (4) is correct, the residuals are not independent because OLS leads to an estimation of $\hat{\boldsymbol{\beta}}$ such that $\sum_{i=1}^n r_i = 0$. However, accounting for the resulting negative correlation between the residuals when creating testing bands would be very difficult, and with a sufficiently large sample size this correlation becomes negligible. Thus, without an a priori reason to believe that the errors may deviate from normality at any particular point in the distribution, it makes sense to use ELL bands in this case. This is very easy to do with **qqconf**, as we show below.

First, we generate data to perform a regression. Here, we will generate each ϵ_i independently from a $t(3)$ distribution.

```
R> set.seed(20)
R> n <- 100
R> x <- runif(n)
R> eta <- rt(n, df = 3)
R> y <- x + eta
```

Then, we fit a regression with the simulated data

```
R> reg <- lm(y ~ x)
```

With **base**-R functionality, we can can create a qqplot using the function `qqnorm` and `qqline`.

```
R> qqnorm(reg$residuals)
R> qqline(reg$residuals)
```

**Normal Q−Q Plot**

Sample Quantiles

Theoretical Quantiles

Figure 5: QQ Plot for Regression Residuals with Base R Functionality

Figure 6: QQ Plot for Regression Residuals with ELL Bounds

Clearly there's some indication of deviation in this plot, but it can be hard to tell how significant the deviation is without confidence bounds. With our function `qq_conf_plot`, it is much easier to be able to judge the significance of the deviation.

```
R> library(qqconf)
R> qq_conf_plot(
+    obs = reg$residuals,
+    points_params = list(col="blue", pch = 20, cex = .5)
+  )
```

Here, we can clearly see that both the left and right tail of the residuals goes beyond the normal testing bounds, giving strong evidence that the errors were not generated from a normal distribution.

This example also highlights an instance in which the ELL method is preferable to KS. Because KS is much more sensitive to deviations in the center of the distribution than it is to deviations in the tails of the distribution, it does not yeild a rejection of the null hypothesis in this case (Figure 7). We can generate KS bounds by simply setting the `method` arguments to `"ks"` in the `qq_conf_plot` function.

```
R> library(qqconf)
```

Figure 7: QQ Plot for Regression Residuals with KS Bounds

```
R> qq_conf_plot(
+     obs = reg$residuals,
+     method = "ks",
+     points_params = list(col="blue", pch = 20, cex = .5)
+   )
```

4.3. Multiple Testing in Genome Wide Association Studies

Applicability of Equal Local Levels to Genome Wide Association Studies

In genome wide association studies (GWAS), we are interested in identifying the genetic variants that are responsible for biological traits. Often, we test a large number of single-nucleotide polymorphisms (SNPs), to see if any of them are strongly correlated with some discrete or continuous trait. Often, it is useful to test the following global null hypothesis:

H_0: All genome-wide SNPs in some functional category have correlation 0 with the trait of interest.

vs.

H_A: At least one SNP in some functional category has non-zero correlation with the trait of interest.

Typically, when testing this global null, we apply some test to each of n SNPs, yielding n p-values. While the SNPs have a local correlation structure, this has only a weak effect at a genome-wide scale. Thus, the local levels procedure still provides a valuable visualization tool to assess the extent and type of deviation from the null, even though it is not exact. Particularly in genetics, if the p-values deviate from the null, it can be very useful to view graphically how they deviate. For instance, if a few tests yield very small p-values but the p-values from the rest of the tests looks relatively uniform, this indicates that the trait of interest is likely driven by a few mutations. If, however, there are some small deviations throughout the p-value distribution, this indicates that trait of interest is likely driven by a number of genes that all play some small part in a complex biological process, or potentially by some confounding variables that are not controlled for.

If a very specific distribution for H_A is suspected, then it makes sense to use a test that has high power under this alternative, which may not be our local levels test. However, in the general case when little is known about the alternative distribution, using local levels is desirable because it is agnostic to the space of alternative distributions. Moreover, since small p-values are often of great interest in GWAS because they can indicate the main biological drivers of a trait, the local levels test is far superior to the KS test due to its comparative tail sensitivity.

Appliction of Equal Local Levels to Creutzfeld-Jakob Disease

We downloaded the p-values from a GWAS of Creutzfeld-Jakob disease in a sample of 4110 cases and 13,569 controls (Jones, Hummerich, Viré, Uphill, Dimitriadis, Speedy, Campbell, Norsworthy, Quinn, Whitfield *et al.* 2020). Test of association between risk for the disease and genetic variants was done at 6,314,492 SNPs. Major genetic risk loci were found on chromosomes 1 and 20. Here, we remove those chromosomes from the results in order to display the p-values from tests that are not obviously from a non-uniform distribution. For display purposes, we subsample the remaining SNPs to 10,000 approximately evenly spaced SNPs. Subsampling helps ensure that correlations between SNPs that are located near each other is minimized.

```
R> cjd_df <- read.table("data/cjd_sample.txt", header = TRUE)
```

We then make QQ plots of these 10,000 tests with 95% testing bands. First, we make a standard QQ plot of observed quantiles versus expected quantiles (Figure 8).

```
R> library(qqconf)
R> qq_conf_plot(
+    obs = cjd_df[,3],
+    distribution = qunif,
+    points_params = list(pch = 21, cex = 0.2)
+    )
```

Unfortunately, this plot results in deviations from the expected line that are difficult to see. To aid visualization, we also plot the deviation of the observed quantiles from the

Figure 8: Standard QQ Plot for CJD GWAS p-values

Figure 9: QQ Plot With y-axis as Observed - Expected Quantile for CJD GWAS p-values

expected quantiles versus the expected quantiles (Figure 9). This can be done easily by setting the `difference` parameter to `TRUE` in `qq_conf_plot`.

```
R> qq_conf_plot(
+    obs = cjd_df[,3],
+    distribution = qunif,
+    points_params = list(pch = 21, cex = 0.2),
+    difference = TRUE
+    )
```

This plot is much easier to interpret, but note that tests in which there is an association will result in small p-values, which will be difficult to see when the plot axes are on the 0 to 1 scale. To see the small p-values we repeat the plots but with the axes on the -log10 scale (Figures 10 and 11). This can be done by setting the `log10` parameter of `qqconf` to `TRUE`

```
R> qq_conf_plot(
+    obs = cjd_df[,3],
+    distribution = qunif,
+    points_params = list(pch = 21, cex = 0.2),
```

Figure 10: QQ Plot With y-axis as -$\log_{10}$(Observed Quantile) for CJD GWAS p-values

```
+     log10 = TRUE
+     )


R> qq_conf_plot(
+     obs = cjd_df[,3],
+     distribution = qunif,
+     points_params = list(pch = 21, cex = 0.2),
+     difference = TRUE,
+     log10 = TRUE,
+     ylim = c(-0.2, 1.1)
+     )
```

From the QQ plots, we can see that there is an excess of small p-values, indicating that the test statistics do not follow the expected null distribution.

5. Discussion

The local levels test using an exact recursive calculation allows for the creation of relatively fast and alternative-agnostic simultaneous testing bands for QQ plots. Of course,

Figure 11: QQ Plot With y-axis as -$\log_{10}$(Observed Quantile) + $\log_{10}$(Expected Quantile) for CJD GWAS p-values

our method is limited by the assumptions that (1) all of the tests are independent and that (2) all of the parameters of the null distribution are known. Currently, we are not aware of methods to handle the dependent case that allow for bands on QQ plots. To handle the issue of parameter uncertainty, Aldor-Noiman has adapted their method such that it will yield exact results when some parameters of the null distribution are unknown. However, if the dataset is very large, Aldor-Noiman's method is prohibitively slow because it uses simulation. In this case, we suggest that our method be used because it is much faster, and with a sufficiently large dataset the uncertainty in the parameters will become negligible under standard regularity conditions. Methods to handle parameter uncertainty with an exact calculation (as opposed to simulation) have been discussed in the context of the normal distribution (Rosenkrantz 2000), but a general method towards this end has not been developed, and would be an important next step.

Our software is available on CRAN. In addition to the QQ plot functionality, we've implemented PP plots using the same methodology and a very similar interface. We've also provided functionality to directly calculate ELL bounds given a sample size n and a global level α. For more information on these functions, please see the code documentation.

References

Aldor-Noiman S, Brown LD, Buja A, Rolke W, Stine RA (2013). "The power to see: A new graphical test of normality." *The American Statistician*, **67**(4), 249–260.

Almeida A, Loy A, Hofmann H (2017). *qqplotr: Quantile-Quantile Plot Extensions for 'ggplot2'*. R package version 0.0.3 initially funded by Google Summer of Code 2017, URL `https://github.com/aloy/qqplotr`.

Berk RH, Jones DH (1979). "Goodness-of-fit test statistics that dominate the Kolmogorov statistics." *Zeitschrift für Wahrscheinlichkeitstheorie und verwandte Gebiete*, **47**(1), 47–59.

Gontscharuk V, Finner H (2017). "Asymptotics of goodness-of-fit tests based on minimum p-value statistics." *Communications in Statistics - Theory and Methods*, **46**(5), 2332–2342. `doi:10.1080/03610926.2015.1041985`. `https://doi.org/10.1080/03610926.2015.1041985`, URL `https://doi.org/10.1080/03610926.2015.1041985`.

Gontscharuk V, Landwehr S, Finner H (2015). "The intermediates take it all: Asymptotics of higher criticism statistics and a powerful alternative based on equal local levels." *Biometrical Journal*, **57**(1), 159–180.

Gontscharuk V, Landwehr S, Finner H, *et al.* (2016). "Goodness of fit tests in terms of local levels with special emphasis on higher criticism tests." *Bernoulli*, **22**(3), 1331–1363.

Jones E, Hummerich H, Viré E, Uphill J, Dimitriadis A, Speedy H, Campbell T, Norsworthy P, Quinn L, Whitfield J, *et al.* (2020). "Identification of novel risk loci and

causal insights for sporadic Creutzfeldt-Jakob disease: a genome-wide association study." *The Lancet Neurology*, **19**(10), 840–848.

Kolmogorov A (1941). "Confidence limits for an unknown distribution function." *The annals of mathematical statistics*, **12**(4), 461–463.

Rosenkrantz WA (2000). "Confidence bands for quantile functions: A parametric and graphic alternative for testing goodness of fit." *The American Statistician*, **54**(3), 185–190.

Smirnov NV (1944). "Approximate laws of distribution of random variables from empirical data." *Uspekhi Matematicheskikh Nauk*, (10), 179–206.

Wilk MB, Gnanadesikan R (1968). "Probability Plotting Methods for the Analysis of Data." *Biometrika*, **55**(1), 1–17. doi:10.1093/biomet/55.1.1.

Affiliation:

Eric Weine
Department of Statistics
University of Chicago
924 E 57th St., Chicago, IL
E-mail: ericweine15@gmail.com

Mary Sara McPeek
Department of Statistics
University of Chicago
924 E 57th St., Chicago, IL
E-mail: mcpeek@uchicago.edu
URL: https://www.stat.uchicago.edu/~mcpeek/

Mark Abney
Department of Human Genetics
University of Chicago
920 E 58th St., Chicago, IL
E-mail: abney@uchicago.edu
URL: https://hgen.uchicago.edu/program/faculty/mark-abney